

# FaceForge Core - MVP Sprint Sequence

---

Below is a sprint sequence that gets us to a shippable **FaceForge Core MVP** without any single sprint turning into a swamp. Each sprint is scoped so you can close it cleanly, tag a milestone, and start the next session with a fresh Copilot thread.

---

## SUMMARY CONCEPTS

### MVP Definition of "Done"

A non-developer can install/run **FaceForge Desktop**, pick `FACEFORGE_HOME`, and then:

- open the local web UI (served by Core)
- create/edit entities
- upload assets, link them to entities, download them (streaming + resume/range)
- run at least one real job (bulk import) and view job logs
- enable/disable a “plugin” (manifest discovery + config surface; plugin compute not required)
- everything binds to localhost by default with a per-install token

(Aligned with the v0.2.9 spec.)

### Optional “buffer” sprints (only if you feel pain)

- **Perf & indexing pass:** SQLite indexes for common list queries, generated columns for frequently filtered JSON fields
- **Backup/export v0:** simple metadata export + asset listing (even if FFBackup becomes a plugin later)
- **API polish:** bulk-upsert endpoint and conflict rules, if you need it early

## Next Steps

If you want to run this like a true sprint board, the cleanest next move is to turn each sprint into an epic, then break deliverables into 5–15 bite-sized issues each (small enough that Copilot doesn’t melt down mid-session).

---

## THE SPRINT SEQUENCE

---

### Sprint 0 — Repo + “boring core” scaffolding

- **Objective:** Make the project easy to work on before you write real features.
- **Deliverables**
  - Monorepo structure (or clean multi-repo) with:
    - `core/` (FastAPI service)
    - `desktop/` (Tauri)

- `docs/` (dev + user docs)
    - Dev bootstrap scripts (one command to run Core in dev, one for UI dev if separate)
    - Formatting/linting/test harness wired (pre-commit or equivalent)
    - CI pipeline: lint + unit tests + build artifacts (at least Core)
  - **Acceptance**
    - Fresh clone → one documented command starts Core + a placeholder `/healthz`
    - CI runs on PRs without manual babysitting
- 

## Sprint 1 — Config + filesystem layout contract

- **Objective:** Lock in how the app finds its home and writes files.
  - **Deliverables**
    - `FACEFORGE_HOME` resolution rules + defaults
    - Directory creation rules (`db/`, `s3/`, `logs/`, `run/`, `config/`, `plugins/`)
    - Core config file format (JSON suggested in spec) + loader + validation
    - Runtime “ports.json” convention (even if Desktop owns it later)
  - **Acceptance**
    - Core starts with only `FACEFORGE_HOME` set and creates required subfolders
    - Config changes require no code changes (just file edits)
- 

## Sprint 2 — SQLite schema v1 + ID strategy

- **Objective:** Build the relational spine cleanly before endpoints pile up.
- **Deliverables**
  - SQLite schema + migrations strategy (simple migration runner is fine)
  - Tables for:
    - `entities`, `assets`, `entity_assets` (many-to-many)
    - `relationships`
    - `jobs`, `job_logs`
    - `field_definitions`
    - `plugin_registry` (enabled/config/version snapshot)
  - SHA-256 ID utilities:
    - entity IDs generated at creation time
    - asset IDs + `content_hash` computed from bytes
  - Soft delete strategy decided (and implemented where relevant)
- **Acceptance**
  - “Create entity” + “create asset record” can be done via a tiny internal script without API

- Migration runner works from blank → latest
- 

## Sprint 3 — Core API skeleton + auth defaults

- **Objective:** A real API surface that won't need rework later.
  - **Deliverables**
    - FastAPI app structure with versioned routing (`/v1/...`)
    - Token auth middleware (per-install token stored in config; localhost by default)
    - Standard response envelope + error model (consistent errors matter)
    - `/docs` and `/redoc` enabled and accurate
  - **Acceptance**
    - A request without token fails (except health endpoints you explicitly exempt)
    - OpenAPI renders and matches real routes
- 

## Sprint 4 — Entities CRUD + list/search primitives

- **Objective:** The simplest “real feature” end-to-end.
  - **Deliverables**
    - Endpoints:
      - `GET /v1/entities`
      - `POST /v1/entities`
      - `GET/PATCH/DELETE /v1/entities/{entity_id}`
    - Minimal filtering/paging/sorting (don't overbuild search yet)
    - Entity model supports: `display_name`, `aliases`, `tags`, `fields` JSON, timestamps
  - **Acceptance**
    - You can create 100 entities and page through them
    - Patch updates only touched fields; timestamps behave sensibly
- 

## Sprint 5 — Assets v1: upload, metadata, link/unlink, download (filesystem provider)

- **Objective:** Make file ingest + retrieval rock-solid before adding S3.
- **Deliverables**
  - Storage abstraction interface + **local filesystem provider** implementation
  - `POST /v1/assets/upload` (multipart; optional companion `_meta.json`)
  - `GET /v1/assets/{asset_id}` metadata
  - Link/unlink endpoints to entities
  - `GET /v1/assets/{asset_id}/download`:
    - streaming response
    - HTTP range support (resume-friendly)

- **IMPORTANT:** Ingest metadata extraction via [exiftool](#)

- Embed [exiftool](#) binary (avoid Perl requirements) - install the correct edition per OS (this should be incorporated into the install process later)
- Run on asset upload to extract metadata
  - Store in asset record under [metadata](#) JSON field
  - Spawn parallel process to run [exiftool](#) if needed (be sure to log events properly)
  - EXCEPTIONS (Files matching these patterns should skip exiftool processing):
    - [\\_\(meta|directorymeta\)\.json\\$](#)
    - [\.\(cover|thumb|thumb\(s|db|index|nail\)\)\\$](#)
    - [^\(thumb|thumb\(s|db|index|nail\)\)\.db\\$](#)
    - [\.\(csv|html?|json|tsv|xml\)\\$](#)
- Send parameters to exiftool via parameter file: [exiftool -@ "\\$ArgsFile" 2> \\$null](#) (in PowerShell, for example)
- Use these exact parameters in the [\\$ArgsFile](#):

```
-quiet -extractEmbedded3 -scanForXMP -unknown2 -json -G3:1 -struct
-b -ignoreMinorErrors -charset filename=utf8 -api requestall=3 -
api largefilesupport=1 --
```

- Parse the output and **REMOVE** the following keys from the JSON before storing:

- [ExifTool:ExifToolVersion](#)
- [ExifTool:FileSequence](#)
- [ExifTool:NewGUID](#)
- [System:BaseName](#)
- [System:Directory](#)
- [System:FileBlockCount](#)
- [System:FileBlockSize](#)
- [System:FileDeviceID](#)
- [System:FileDeviceNumber](#)
- [System:FileGroupID](#)
- [System:FileHardLinks](#)
- [System:FileInodeNumber](#)
- [System:FileName](#)
- [System:FilePath](#)
- [System:FilePermissions](#)
- [System:FileUserID](#)

- The final resulting JSON data ([\\$ExifToolOutput](#)) must be validated as proper JSON and being not-empty and nested into the following structure before storing:

```
{
  "Source": "ExifTool",
  "Type": "JsonMetadata",
  "Name": $null,
  "NameHashes": $null,
```

```
    "Data": {ExifToolOutput}  
}
```

- **Acceptance**

- Upload a 2–5GB file and download it reliably
  - Range requests work (spot-check with curl / a download manager)
  - Uploaded asset gets metadata extracted and stored correctly (with proper error handling/logging/validation)
- 

## Sprint 6 — SeaweedFS provider + “default local S3” wiring

- **Objective:** Bring the spec’s default object storage online without dragging Desktop into it yet.
- **Deliverables**

- SeaweedFS “managed binary” contract defined (paths, args, health checks)
- S3-compatible provider implementation via SeaweedFS endpoint
- Storage routing rules (by asset **kind** or size threshold) working
- Core can operate even if SeaweedFS is disabled (fallback to filesystem provider)

- **Acceptance**

- With SeaweedFS running: uploads land in the S3 backend
  - With SeaweedFS off: uploads still work (filesystem)
- 

## Sprint 7 — Descriptors + Field Definitions (admin)

- **Objective:** Get the “flexible schema without migrations” concept working.
- **Deliverables**

- Field definition endpoints:
  - **GET/POST/PATCH/DELETE** /v1/admin/field-defs...
- Descriptor endpoints:
  - **GET/POST** /v1/entities/{entity\_id}/descriptors
  - **PATCH/DELETE** /v1/descriptors/{descriptor\_id}
- Validation driven by field definitions (at least type + required + regex/options)

- **Acceptance**

- Add a new field definition → UI/API accepts it immediately
  - Invalid descriptor values are rejected with clear errors
- 

## Sprint 8 — Relationships + relation type suggestion

- **Objective:** Relationship metadata, without graph visualization.
- **Deliverables**

- `POST /v1/relationships` and `GET /v1/relationships?entity_id=...`
- `DELETE /v1/relationships/{relationship_id}`
- `GET /v1/relationship-types?query=...` backed by:
  - existing relationship types in DB
  - a small built-in seed list (so first-time UX isn't empty)

- **Acceptance**

- Create relationship types organically; suggestion endpoint surfaces them
  - Relationship records round-trip correctly and are queryable
- 

## Sprint 9 — Jobs + structured logs + one “real” job

- **Objective:** A durable job model that plugins can piggyback on later.

- **Deliverables**

- `POST /v1/jobs`, `GET /v1/jobs/{job_id}`, `GET /v1/jobs/{job_id}/log`, `POST /v1/jobs/{job_id}/cancel`
- Append-only structured logging with timestamps + levels
- Progress reporting fields (percent + step name)
- Implement **one Core job** end-to-end:
  - `assets/bulk-import` as a job (directory scan + sidecar `_meta.json`)

- **Acceptance**

- Start bulk import → watch progress/logs live → job completes
  - Cancellation behaves predictably (stops work, marks state)
- 

## Sprint 10 — Plugin foundations v1 (discovery + registry + config surface)

- **Objective:** “Plugins exist” without building plugin compute yet.

- **Deliverables**

- Plugin manifest format (`plugin.json`) and loader
- Discovery from `#{FACEFORGE_HOME}/plugins`
- Registry persisted in DB:
  - enabled/disabled
  - stored config blob (validated against JSON schema if provided)
- API:
  - `GET /v1/plugins`
  - enable/disable
  - `GET/PUT /v1/plugins/{id}/config`
- Namespacing rules established (`/v1/plugins/<id>/...` reserved)

- **Acceptance**

- Drop a manifest into plugins dir → it appears in API/UI

- Enable/disable state persists across restarts
- 

## Sprint 11 — Core Web UI MVP (Core-only features)

- **Objective:** A usable UI that covers everything Core can do (no “backend-only” gaps).

- **Deliverables**

- Pages:
  - Entities list (table + gallery toggle)
  - Entity detail with tabs: Overview / Descriptors / Attachments / Relationships
  - Jobs page: list + detail log view
  - Plugins page: list + enable/disable + config form rendering (basic)
- Attachments UX:
  - upload
  - link/unlink
  - download
- Lightweight styling, fast load, no runtime Node dependency

- **Acceptance**

- A non-technical user can do the whole “create entity → upload asset → link → download” loop without touching API tools
  - Jobs are viewable and understandable from UI
- 

## Sprint 12 — Desktop shell MVP (Tauri orchestrator + tray UX)

- **Objective:** Make it run like a real app, not a repo full of scripts.

- **Deliverables**

- First-run wizard:
  - choose `FACEFORGE_HOME`
  - choose/auto-pick Core port (and SeaweedFS port if needed)
- Process orchestration:
  - start/stop Core service
  - start/stop SeaweedFS (if enabled)
  - basic health monitoring + restart strategy
- Tray behavior:
  - close to tray (not exit)
  - menu: Open UI, Status, Logs, Stop, Restart, Exit
  - on exit: prompt to stop services or leave running

- **Acceptance**

- Double-click app → wizard → “Open UI” works
  - Closing window leaves tray running; Exit performs expected shutdown behavior
-

## Sprint 13 — Packaging + release hardening

- **Objective:** Turn “works on my machine” into “ship it.”
  - **Deliverables**
    - Packaging plan implemented:
      - Core service shipped as frozen binary or embedded Python runtime (pick one and commit)
      - SeaweedFS binaries bundled or fetched in a controlled way
    - Installer / portable bundle per OS (Windows first if that’s your priority)
    - Logging + crash reports in  `${FACEFORGE_HOME}/logs`
    - Security defaults verified:
      - localhost bind by default
      - token required
      - config/secrets not accidentally committed
    - “First-run screenshots + quickstart” user docs
  - **Acceptance**
    - Clean machine install → user completes wizard → app works without dev tooling
    - Basic smoke test checklist passes reliably
-